

```
[52] def customer_recomendation(customer_id):  
      if customer_id not in df_output.index:  
          print('Customer not found.')  
          return customer_id  
      return df_output.loc[customer_id]
```

```
[53] customer_recomendation(4)
```

```
➞ recommendedProducts    2|1|36|13|216|61|20|33|25|157  
   Name: 4, dtype: object
```

```
[54] customer_recomendation(21)
```

```
➞ recommendedProducts    38|36|48|79|1|2|15|13|25|20  
   Name: 21, dtype: object
```

▼ 7. Model Evaluation

For evaluating recommendation engines, we can use the concept of precision-recall.

- RMSE (Root Mean Squared Errors)
 - Measures the error of predicted values
 - Lesser the RMSE value, better the recommendations
- Recall
 - What percentage of products that a user buys are actually recommended?
 - If a customer buys 5 products and the recommendation decided to show 3 of them, then the recall is 0.6
- Precision
 - Out of all the recommended items, how many the user actually liked?
 - If 5 products were recommended to the customer out of which he buys 4 of them, then precision is 0.8

Lets compare all the models we have built based on precision-recall characteristics:

8.1. Evaluation summary

- Based on RMSE

1. Popularity on purchase counts: 1.1111750034210488
2. Cosine similarity on purchase counts: 1.9230643981653215
3. Pearson similarity on purchase counts: 1.9231102838192284

4. Popularity on purchase dummy: 0.9697374361161925
5. Cosine similarity on purchase dummy: 0.9697509978436404
6. Pearson similarity on purchase dummy: 0.9697745320187097

7. Popularity on scaled purchase counts: 0.16230660626840343
8. Cosine similarity on scaled purchase counts: 0.16229800354111104
9. Pearson similarity on scaled purchase counts: 0.1622982668334026

- Based on Precision and Recall

```
name = 'pearson'
target = 'purchase_count'
pear = model(train_data, name, user_id, item_id, target, users_to_recommend, n_rec, n_display)
```

Preparing data set.
Data has 106868 observations with 23382 users and 300 items.
Data prepared in: 0.154145s

Training model from provided data.
Gathering per-item and per-user statistics.

Elapsed Time (Item Statistics)	% Complete
1.484ms	4.25
24.436ms	100

Setting up lookup tables.
Processing data in one pass using dense lookup tables.

Elapsed Time (Constructing Lookups)	Total % Complete	Items Processed
25.339ms	0	0
77.412ms	100	300

Finalizing lookup tables.
Generating candidate set for working with new users.
Finished training in 0.095617s
recommendations finished on 1000/1000 queries. users per second: 47321.6

customerId	productId	score	rank
1553	248	3.269303671338341	1
1553	132	3.172413793103448	2
1553	37	3.0322580645161303	3
1553	34	2.9947963457127105	4
1553	0	2.9891616797683294	5
1553	3	2.7993184668951274	6
1553	110	2.7272471324963994	7
1553	27	2.69852878164528	8
1553	230	2.673649873986517	9
1553	10	2.6516516516516515	10
20400	248	3.2727272727272725	1
20400	132	3.172413793103448	2
20400	37	3.029654354818407	3
20400	34	2.995901639344264	4
20400	0	2.992079207920795	5
20400	3	2.7986623836689235	6

recommendations finished on 1000/1000 queries. users 1

customerId	productId	score	rank
1553	259	1.0	1
1553	44	1.0	2
1553	298	1.0	3
1553	45	1.0	4
1553	249	1.0	5
1553	106	1.0	6
1553	15	1.0	7
1553	3	1.0	8
1553	49	1.0	9
1553	11	1.0	10
20400	44	1.0	1
20400	298	1.0	2
20400	45	1.0	3
20400	249	1.0	4
20400	106	1.0	5
20400	15	1.0	6
20400	3	1.0	7
20400	49	1.0	8
20400	193	1.0	9
20400	11	1.0	10
19750	44	1.0	1
19750	298	1.0	2
19750	45	1.0	3
19750	249	1.0	4
19750	106	1.0	5
19750	15	1.0	6
19750	3	1.0	7
19750	49	1.0	8
19750	193	1.0	9
19750	11	1.0	10

```
name = 'cosine'
target = 'purchase_count'
cos = model(train_data, name, user_id, item_id, target, users_to_recommend, n_rec, n_di
```

Preparing data set.
Data has 106868 observations with 23382 users and 300 items.
Data prepared in: 0.169002s
Training model from provided data.
Gathering per-item and per-user statistics.

Elapsed Time (Item Statistics)	% Complete
1.202ms	4.25
10.97ms	100

Setting up lookup tables.
Processing data in one pass using dense lookup tables.

Elapsed Time (Constructing Lookups)	Total % Complete	Items Processed
12.926ms	0	0
49.811ms	100	300

Finalizing lookup tables.
Generating candidate set for working with new users.
Finished training in 1.07754s
recommendations finished on 1000/1000 queries. users per second: 33992.8

customerId	productId	score	rank
1553	2	0.07160244882106781	1
1553	35	0.0693587213754654	2
1553	1	0.0691552609205246	3
1553	61	0.06251166760921478	4
1553	21	0.04497094452381134	5
1553	76	0.04215094447135925	6
1553	0	0.03977116942405701	7
1553	269	0.037718966603279114	8
1553	8	0.03595167398452759	9
1553	5	0.03342823684215546	10
20400	122	0.045564889907836914	1
20400	215	0.04000645875930786	2
20400	6	0.03698962926864624	3
20400	1	0.036445021629333496	4
20400	54	0.03542596101760864	5
20400	56	0.034662067890167236	6
20400	179	0.0340539813041687	7

customerId	productId	score	rank
0	248	3.272727272727273	1
0	132	3.1724137931034484	2
0	37	3.032258064516129	3
0	34	2.9959016393442623	4
0	0	2.992079207920792	5
0	3	2.800429184549356	6
0	110	2.7329545454545454	7
0	27	2.699248120300752	8
0	230	2.676258992805755	9
0	10	2.6516516516516515	10
1	248	3.272727272727273	1
1	132	3.1724137931034484	2
1	37	3.032258064516129	3
1	34	2.9959016393442623	4
1	0	2.992079207920792	5
1	3	2.800429184549356	6
1	110	2.7329545454545454	7
1	27	2.699248120300752	8
1	230	2.676258992805755	9
1	10	2.6516516516516515	10
2	248	3.272727272727273	1
2	132	3.1724137931034484	2
2	37	3.032258064516129	3
2	34	2.9959016393442623	4
2	0	2.992079207920792	5
2	3	2.800429184549356	6
2	110	2.7329545454545454	7
2	27	2.699248120300752	8
2	230	2.676258992805755	9
2	10	2.6516516516516515	10

customerid	productid	purchase_count
1750	146	2
2870	76	1
21091	4	2
9881	2	1
10807	189	1
8564	274	1
13417	1	2
25794	228	1
21384	228	2
15982	175	1

3.3. Normalize item values across users

- To do this, we normalize purchase frequency of each item across users by first creating a user-item matrix as follows

```
[14] df_matrix = pd.pivot_table(data, values='purchase_count', index='customerId', columns='productId')
      df_matrix.head()
```

	productId	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26	27
customerId																													
0		NaN	2.0	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	1.0	NaN	NaN	NaN	NaN	NaN	3.0	1.0	NaN	NaN	NaN	NaN	NaN	NaN	NaN
1		NaN	NaN	6.0	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	1.0	NaN	1.0	NaN	NaN
2		NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	3.0	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	1.0	NaN	NaN	NaN	NaN
3		NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	1.0	NaN	NaN
4		NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	2.0	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN

5 rows × 300 columns

```
[ 6] print(customers.shape)
      customers.head()
```

(1000, 1)

customerId	
0	1553
1	20400
2	19750
3	6334
4	27773

```
[ 7] print(transactions.shape)
      transactions.head()
```

(62483, 2)

customerId		products
0	0	20
1	1	2 2 23 68 68 111 29 86 107 152
2	2	111 107 29 11 11 11 33 23
3	3	164 227
4	5	2 2

User Based Collaborative Filtering

The goal of user-based collaborative filtering is to predict the rating of i_n given the previous actions of one user, u_m . In user-based collaborative filtering, we find k candidates of other users that are similar to u_m . To find similarity, we can use

Cosine Similarity:

$$\cos\theta = \frac{A \cdot B}{|A||B|} = \frac{\sum_{i=1}^n A_i B_i}{\sqrt{\sum_{i=1}^n A_i^2} \sqrt{\sum_{i=1}^n B_i^2}}$$

Pearson Correlation:
$$r = \frac{\sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y})}{\sqrt{\sum_{i=1}^n (x_i - \bar{x})^2} \sqrt{\sum_{i=1}^n (y_i - \bar{y})^2}}$$

There is also a method of using K-nearest-neighbors to find the most similar candidate users, defined by

$$p_{a,i} = \bar{r}_a + \frac{\sum_{u \in K} (r_{u,i} - \bar{r}_u) \times w_{a,u}}{\sum_{u \in K} w_{a,u}},$$

where $p_{a,i}$ is the prediction for target user u_m at item i_n , $w_{a,u}$ is the similarity between the users, and K is the neighborhood of similar users.

Item-based Collaborative Filtering

The goal of item-based collaborative filtering is to predict ratings for a specific item i_n based on ratings of similar previous items rated by all possible users.

The rating for target item i_n for active user u_m can be found using

Weighted Average

$$p_{a,i} = \frac{\sum_{j \in K} r_{a,j} w_{i,j}}{\sum_{j \in K} |w_{i,j}|}$$

where K is the neighborhood of most similar items rated by user u_m and $w(i, j)$ is the similarity between items i and j .

Adjusted Cosine Similarity

$$\text{sim}_{ij} = \frac{\sum_{u \in U} (R_{u,i} - \bar{R}_u)(R_{u,j} - \bar{R}_u)}{\sqrt{\sum_{u \in U} (R_{u,i} - \bar{R}_u)^2} \sqrt{\sum_{u \in U} (R_{u,j} - \bar{R}_u)^2}}$$

	i_1	i_2	i_3	i_4	\dots	i_n
u_1	✓		✓			
u_2			✓			✓
u_3				✓		
u_4		✓		✓		
\dots						
u_m			✓			